



scientist.com

The RDKit is getting Rust-y...

Javier Pineda, PhD

11 September 2024



Overview

1. Scientist.com

Company context and structure search requirements

2. The Rust-y RDKit

Low-level Rust bindings for the RDKit C++

3. Cheminée

Using the Rust-y RDKit for a stand-alone structure search application

4. Future Directions



Overview

1. Scientist.com

Company context and structure search requirements

2. The Rust-y RDKit

Low-level Rust bindings for the RDKit C++

3. Cheminée

Using the Rust-y RDKit for a stand-alone structure search application

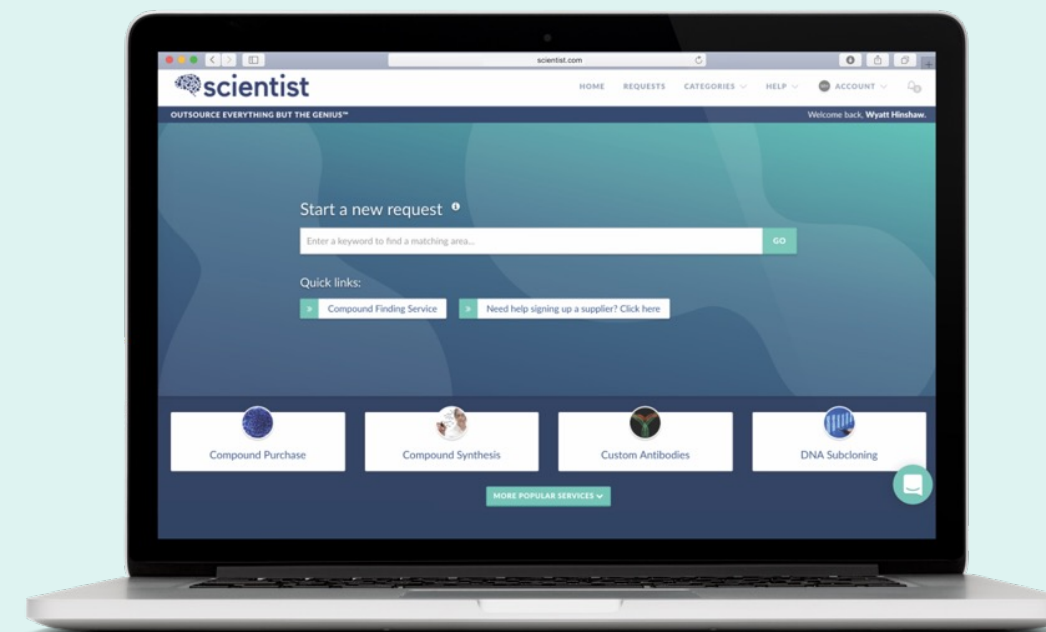
4. Future Directions



Scientist.com is an R&D marketplace for researchers



Large and Small
Research Organizations



Scientist.com Marketplace




Global Suppliers of Life
Science Products &
Services

Human Biological Samples • Compound Synthesis • Antibody Services • Protein Sciences • Toxicology
In Vitro Assays • DNA/RNA Services • Assay Development • Molecular Pharmacology
Bioanalysis • Cell Culture and Characterization • Data Sciences • CMC • Proteomics



Compound searches require structure searches

- ✓ 10+ million compounds from dozens of distributors of compounds
- ✓ Potential for scale up (e.g. more distributors, virtual compounds)
- ✓ We want something simple to spin up locally for our software dev team



1-[3-(1-hydroxyethyl)phenyl]propan-1-one?





Overview

1. Scientist.com

Company context and structure search requirements

2. The Rust-y RDKit

Low-level Rust bindings for the RDKit C++

3. Cheminée

Using the Rust-y RDKit for a stand-alone structure search application

4. Future Directions

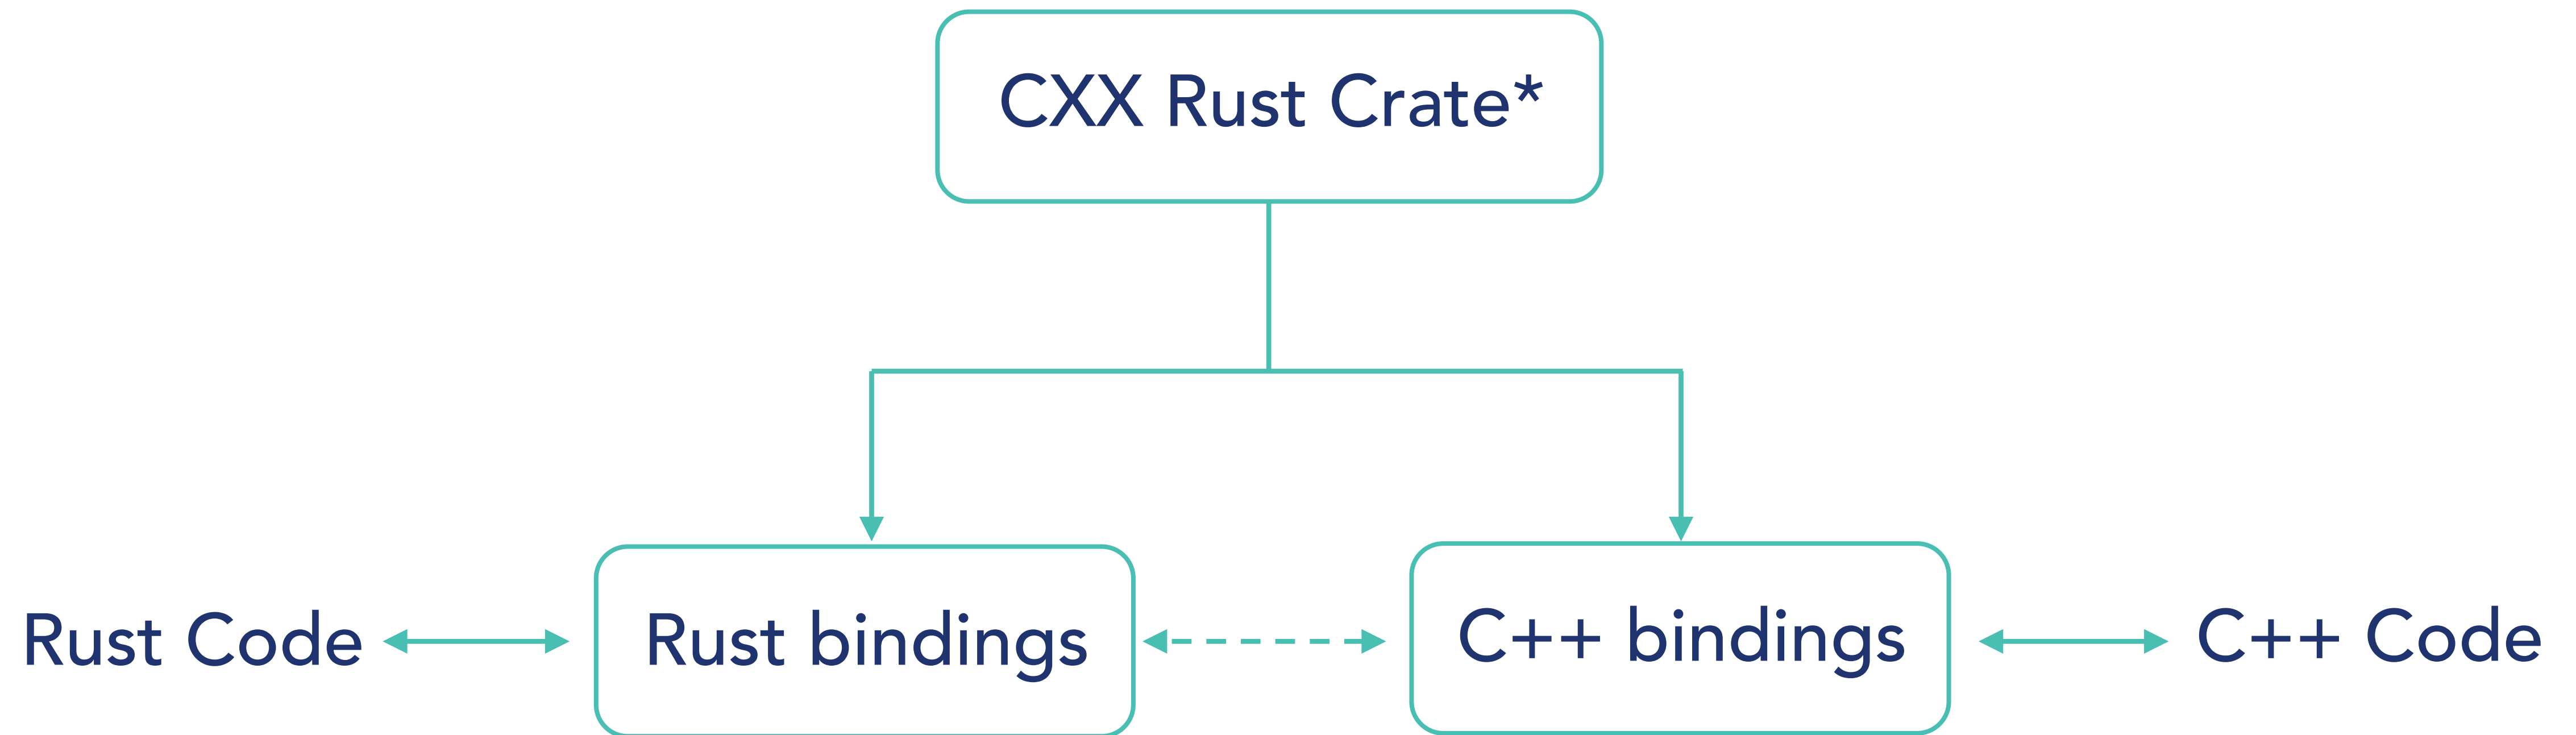


Why Rust?

	C++	Python	Rust
Fast	Yes	No	Yes
Memory Safe	If you code right	Garbage collector	Yes
Thread Safe	If you code right	If you code right	Yes
User Friendly	No	Yes	Kind of



How should we implement the RDKit in Rust?



*<https://cxx.rs>



How is this different from the RDKit CFFI?

RDKit CFFI*

Pros

- Already tied to RDKit
- User friendly
- Agnostic of end language

Cons

- Indirect bindings
- Serialization/deserialization
- Limited access to functionality

Rust CXX

- Directly call C++ code from Rust
- No serialization/deserialization
- Unlimited access to functionality
- Better handling of exceptions

- Limited to Rust
- Not yet implemented for RDKit



Creating Rust bindings can be cumbersome

Higher-level
Rust code

```
// rdkit/src/descriptors.rs
use std::collections::HashMap;
use cxx::SharedPtr;
use crate::ROMol;

pub struct Properties {
    ptr: SharedPtr<rdkit_sys::descriptors_ffi::Properties>,
}

impl Default for Properties {
    fn default() -> Self {
        Properties::new()
    }
}

impl Properties {
    pub fn new() -> Self {
        Self {
            ptr: rdkit_sys::descriptors_ffi::new_properties(),
        }
    }

    pub fn compute_properties(&self, ro_mol: &ROMol) ->
        HashMap<String, f64> {
        let names = rdkit_sys::descriptors_ffi::get_property_
            names(&self.ptr);
        let computed = rdkit_sys::descriptors_ffi::compute_pr
            operties(&self.ptr, &ro_mol.ptr);

        assert!(names.len() != 0);
        assert!(computed.len() == names.len());

        names
            .into_iter()
            .zip(computed.as_slice())
            .map(|(k, v)| (k.to_string(), *v))
            .collect()
    }
}
```

Rust Bindings

```
// rdkit/rdkit-sys/src/bridge/descriptors.rs
#[cxx::bridge(namespace = "RDKit")]
pub mod ffi {
    unsafe extern "C++" {
        include!("wrapper/include/ro_mol.h");
        include!("wrapper/include/descriptors.h");

        pub type ROMol = crate::ro_mol_ffi::ROMol;
        pub type Properties;

        pub fn new_properties() -> SharedPtr<Properties>;
        pub fn get_property_names(
            properties: &SharedPtr<Properties>,
        ) -> UniquePtr<CxxVector<CxxString>>;
        pub fn compute_properties(
            properties: &SharedPtr<Properties>,
            mol: &SharedPtr<ROMol>,
        ) -> UniquePtr<CxxVector<f64>>;
    }
}
```

CXX Bridge*

```
// rdkit/rdkit-sys/wrapper/wrc/descriptors.cc
#include "rust/cxx.h"
#include <GraphMol/Descriptors/Property.h>
#include <GraphMol/ROMol.h>

namespace RDKit {
    using Descriptors::Properties;

    std::shared_ptr<Properties> new_properties() {
        return std::shared_ptr<Properties>(new Properties());
    }

    std::unique_ptr<std::vector<std::string>>
    get_property_names(const std::shared_ptr<Properties>
        &props) {
        std::vector<std::string> names =
            props->getPropertyNames();
        std::vector<std::string> *names_heap = new
            std::vector<std::string>(names);
        return std::unique_ptr<std::vector<
            std::string>>(names_heap);
    }

    std::unique_ptr<std::vector<double>> compute_properties(
        const std::shared_ptr<Properties> &props, const
        std::shared_ptr<ROMol> &mol) {
        std::vector<double> computed =
            props->computeProperties(*mol);
        auto computed_heap = new std::vector<
            double>(computed);
        return std::unique_ptr<std::vector<
            double>>(computed_heap);
    }
}
```

```
// rdkit/rdkit-sys/wrapper/descriptors.h
#pragma once
#include "rust/cxx.h"
#include <GraphMol/Descriptors/Property.h>

namespace RDKit {
    using Descriptors::Properties;

    std::shared_ptr<Properties> new_properties();
    std::unique_ptr<std::vector<std::string>>
    get_property_names(const std::shared_ptr<Properties>
        &props);
    std::unique_ptr<std::vector<double>> compute_properties(
        const std::shared_ptr<Properties> &props, const
        std::shared_ptr<ROMol> &mol);
}
```

C++
Bindings

*In this case, just for chemical descriptor methods




What about automated Rust bindings to RDKit?

- ✓ An attractive idea, considering only two of us are working on this
- ✓ AUTOCXX Rust crate* is already built to automate Rust-to-C++ bindings



Automated bindings!



Did not work reliably

*<https://docs.rs/autocxx>



Back to manual RDKit Rust bindings!

Mol*

from_smiles	atom_with_idx
to_smiles	clean_up
from_molblock	detect_chemistry_problems
to_molblock	update_property_cache
fingerprint	substruct_match
add_hs	descriptors
remove_hs	tautomer_canonicalization

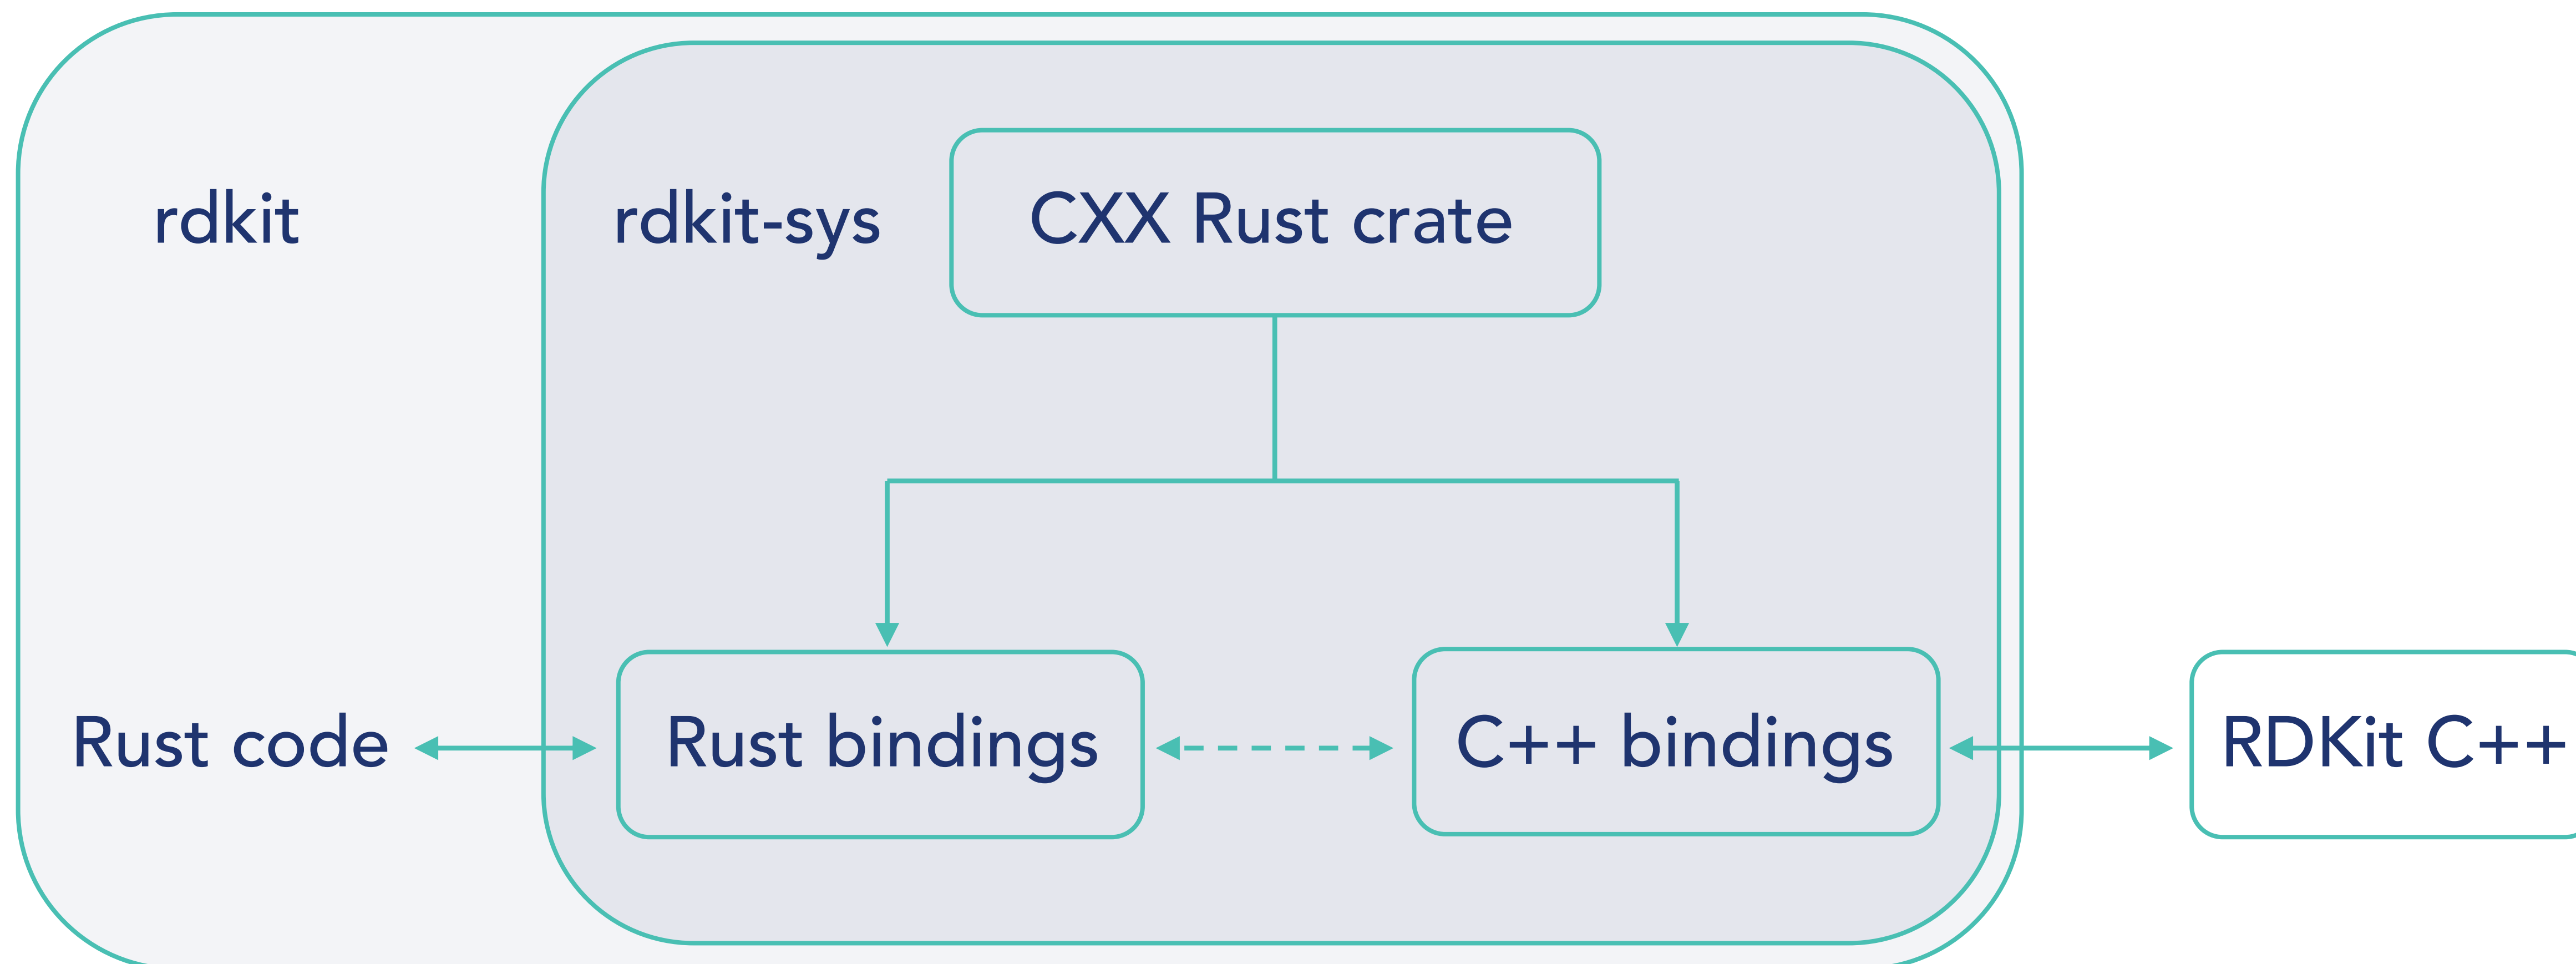
Atom*

get_symbol
is_aromatic
get_formal_charge
hybridization
explicit_hs
total_valence
atomic_num

*A few of the bindings we've created



Say hello to the Rust-y RDKit*



*Repo: <https://github.com/rdkit-rs/rdkit>

Crate: <https://docs.rs/rdkit>



Overview

1. Scientist.com

Company context and structure search requirements

2. The Rust-y RDKit

Low-level Rust bindings for the RDKit C++

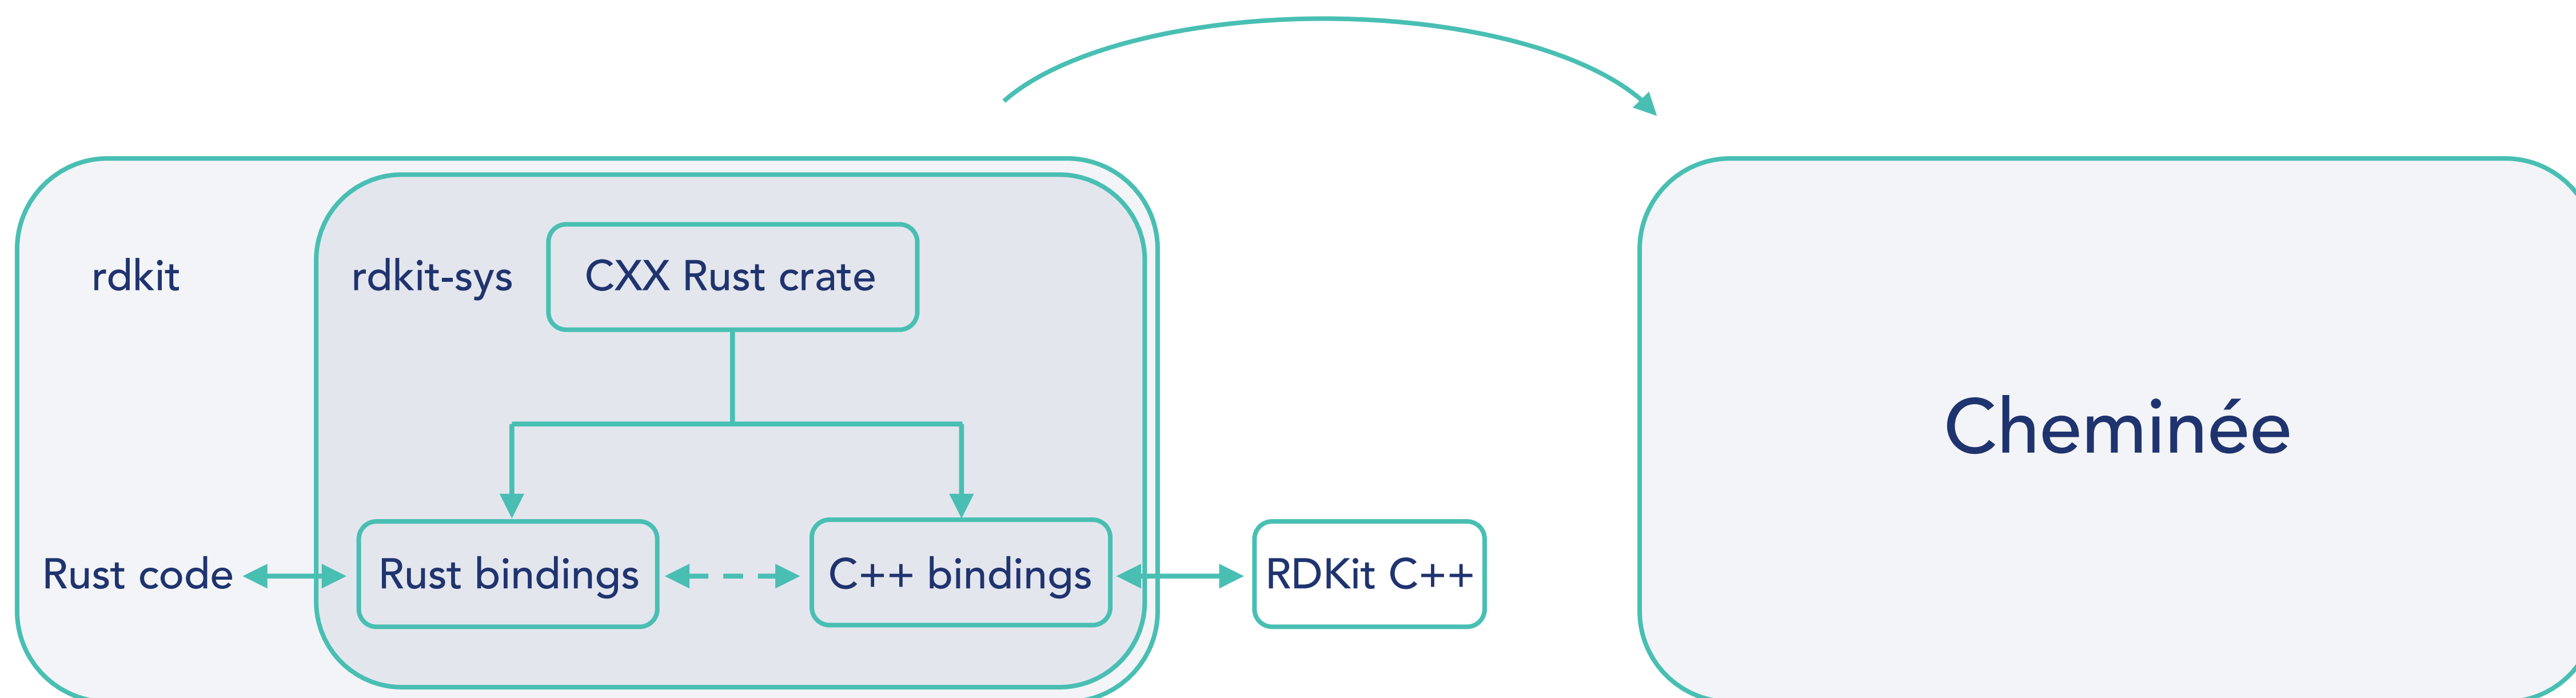
3. Cheminée

Using the Rust-y RDKit for a structure search application

4. Future Directions

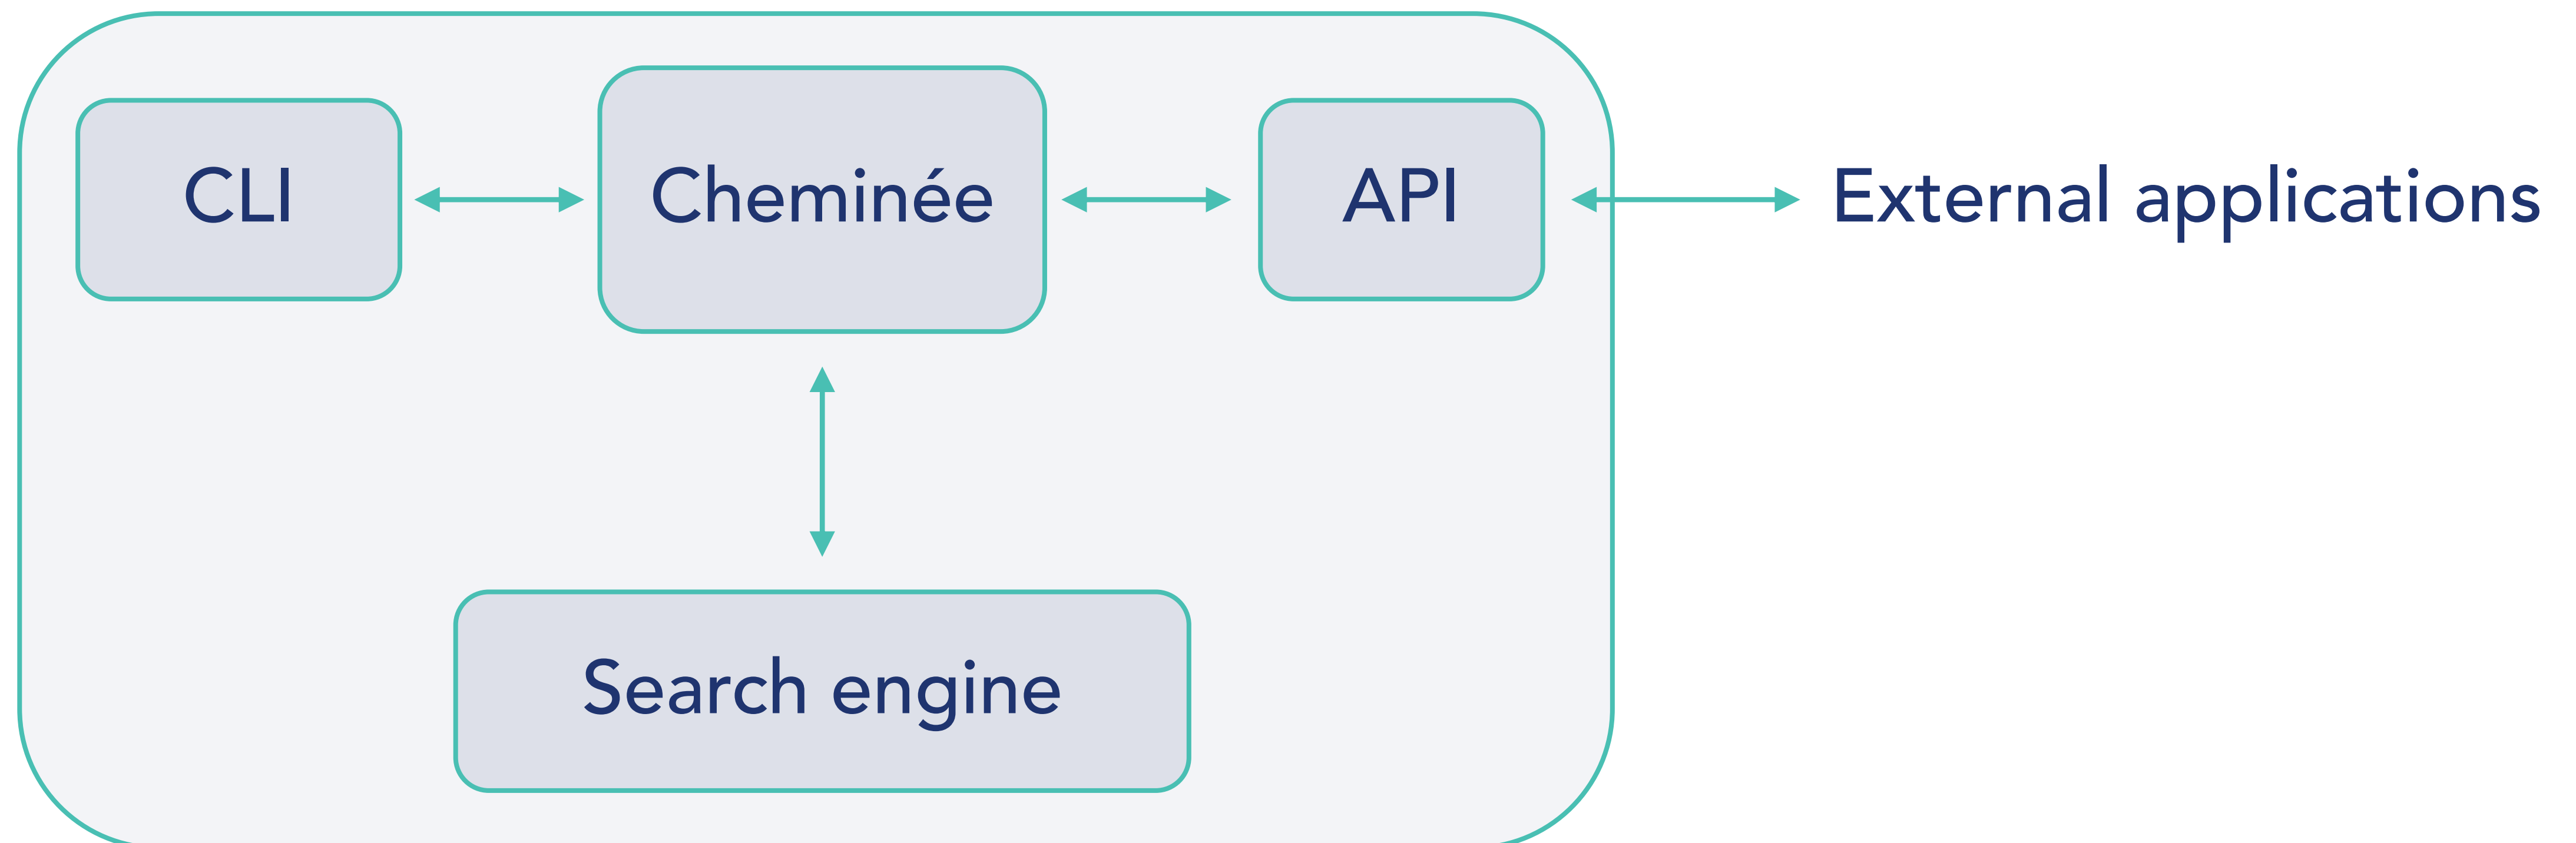


Cheminée is our structure search engine



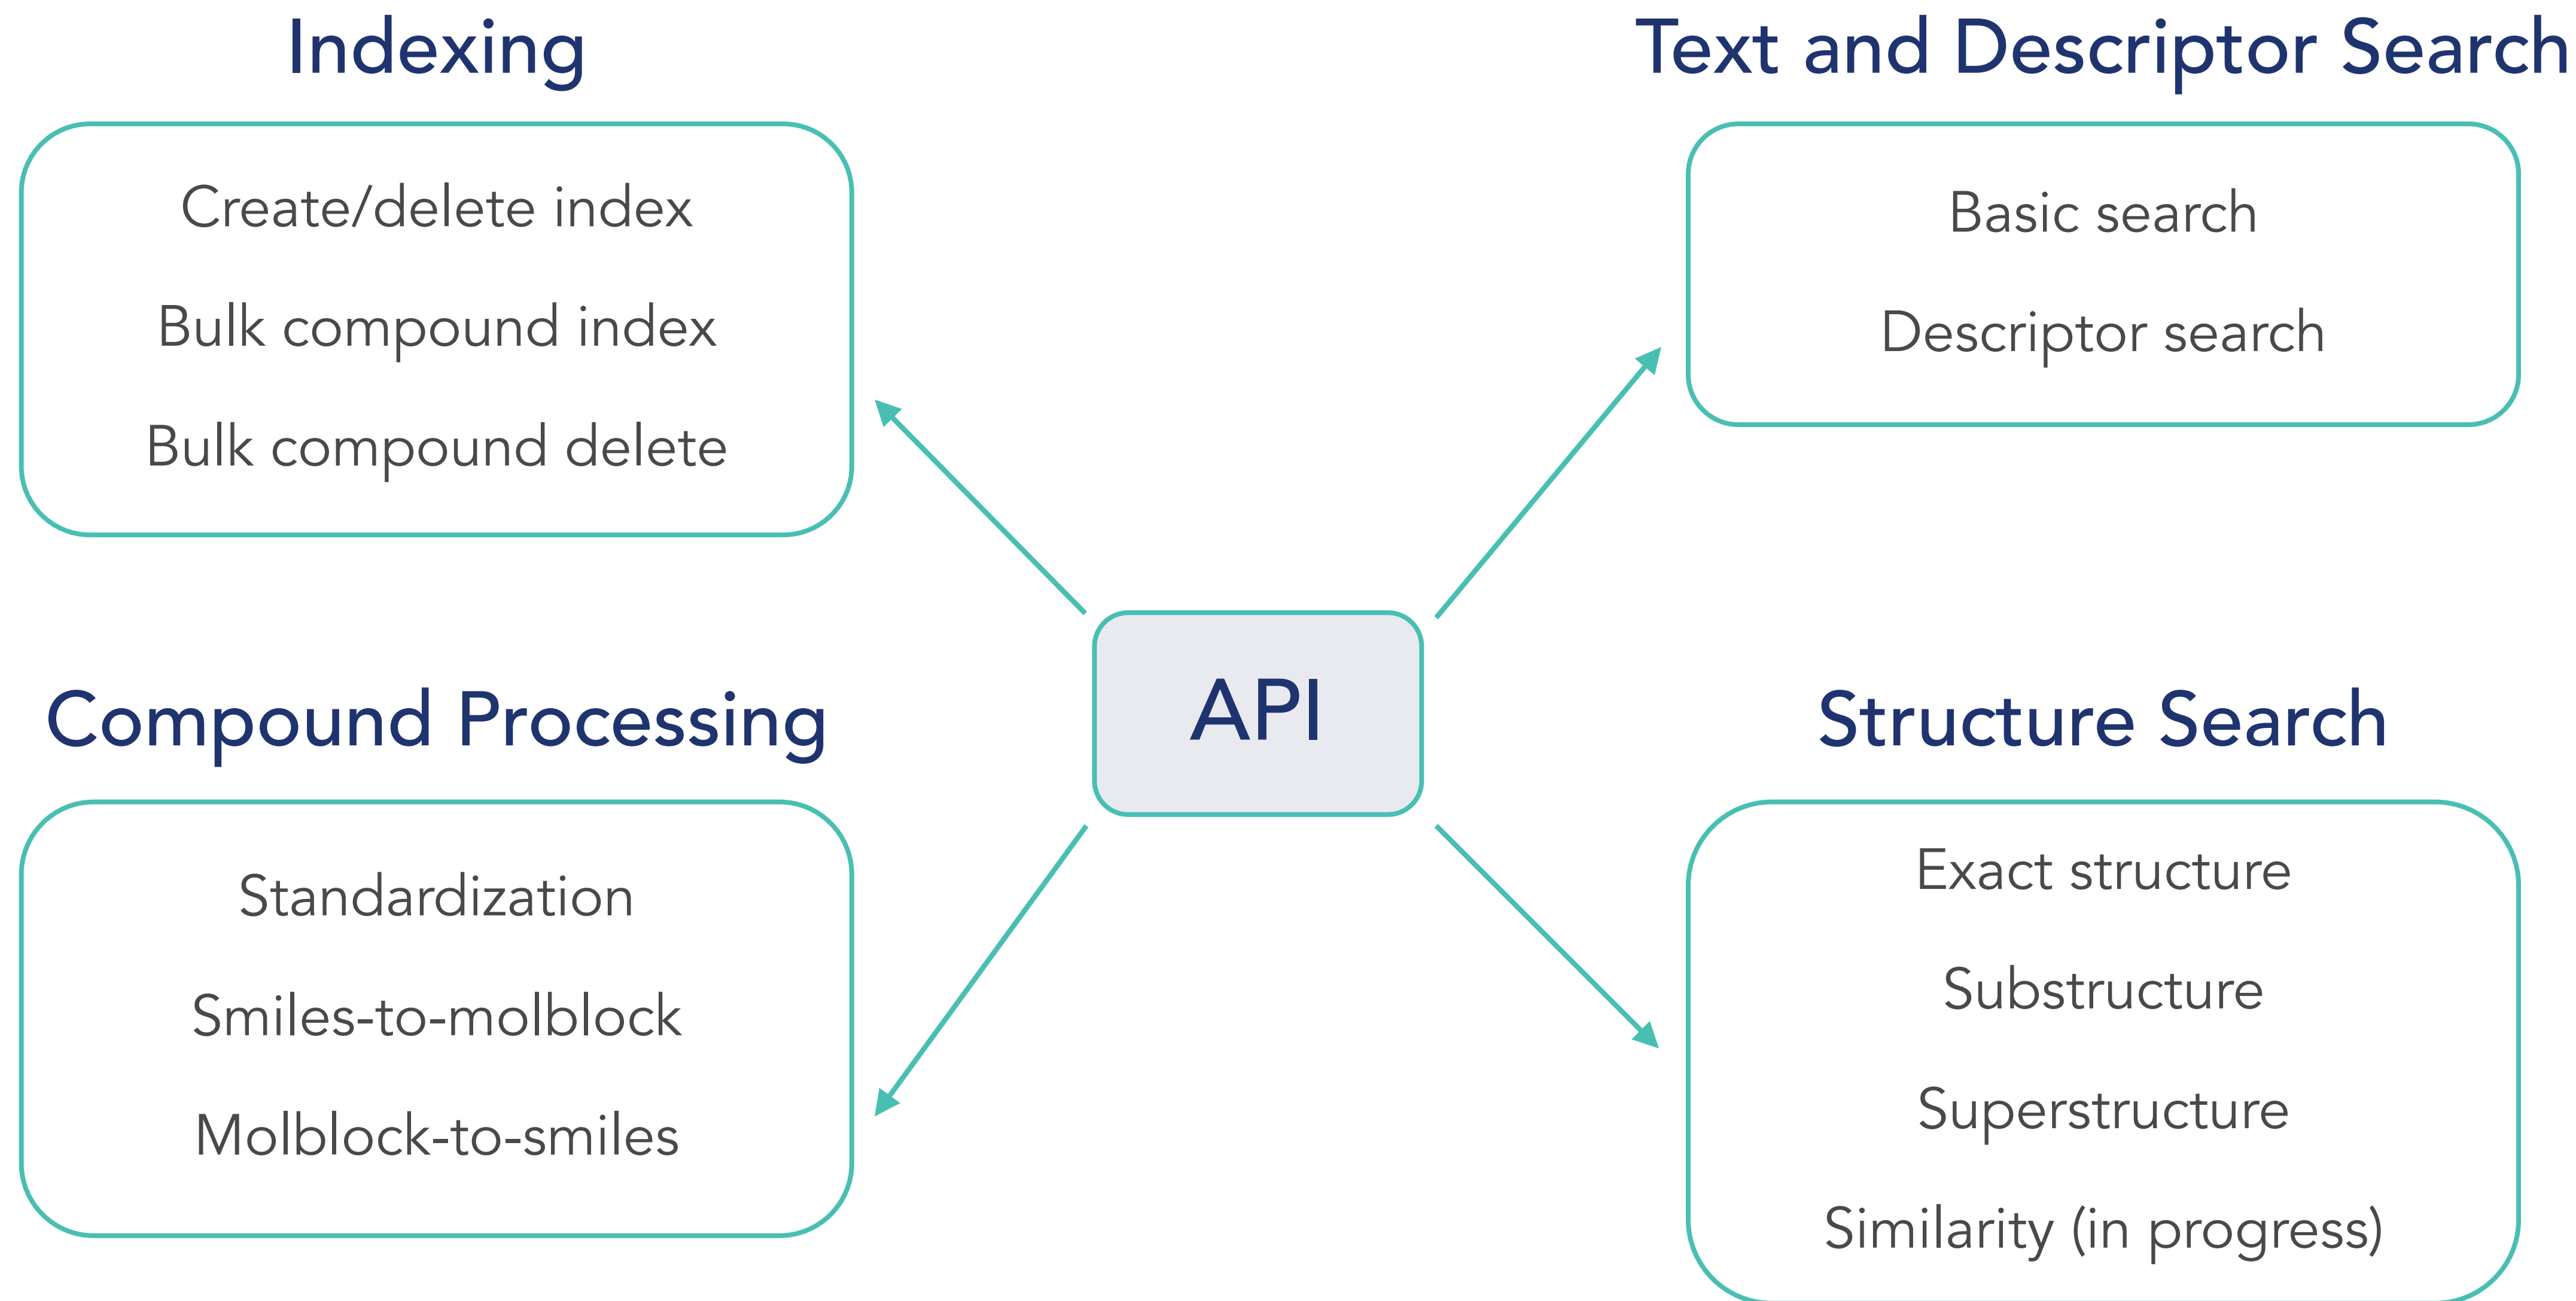


Cheminée has its own CLI, API, and search engine



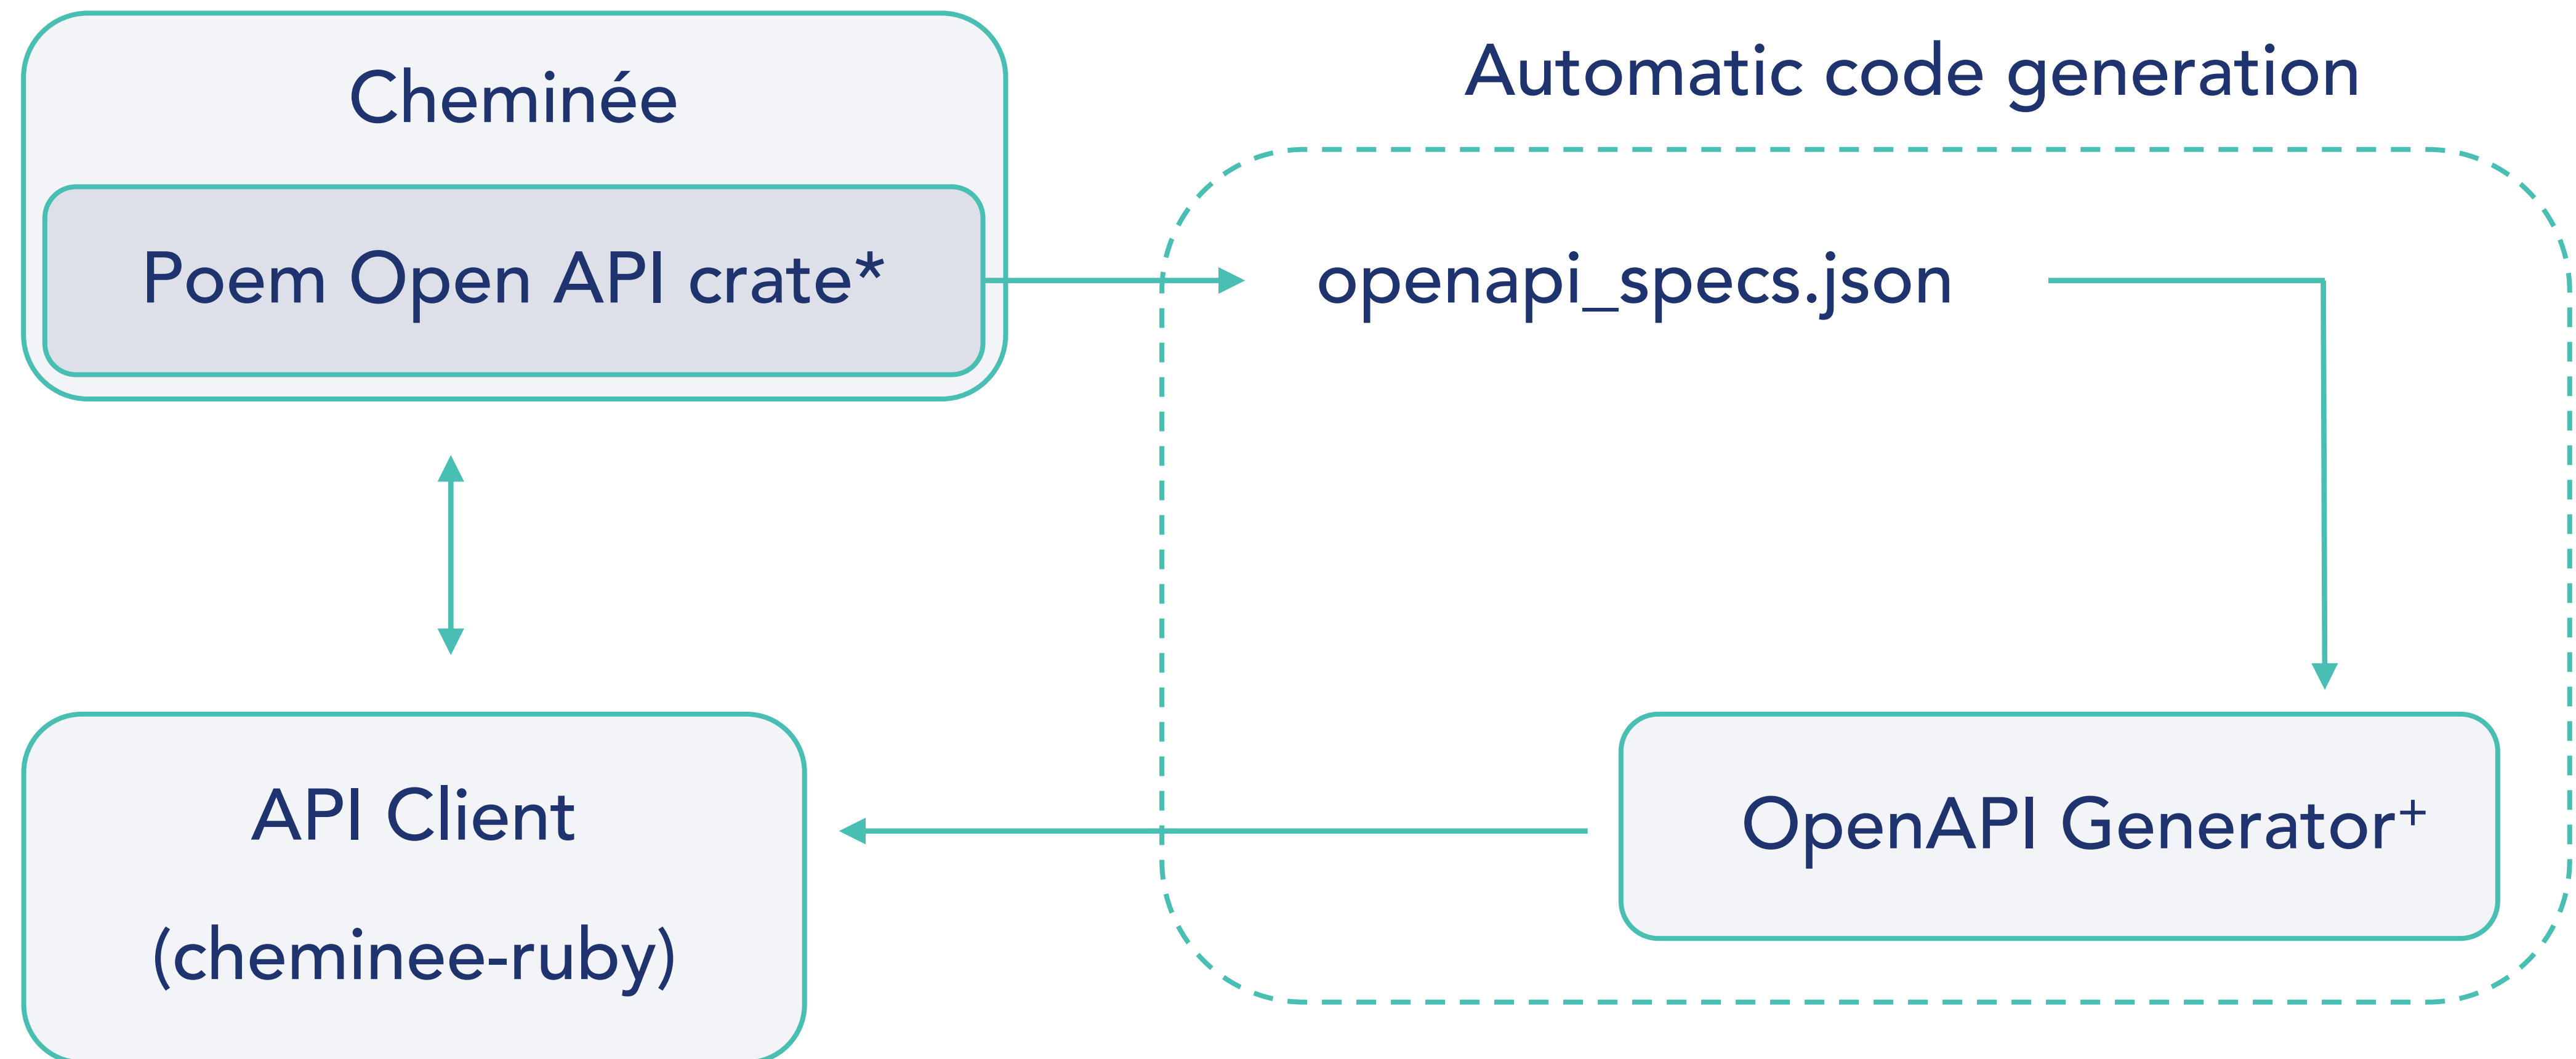


Cheminée has several endpoints





The Cheminée API works smarter, not harder



*<https://docs.rs/poem-openapi>

+<https://github.com/OpenAPITools/openapi-generator>



Cheminée does not use the RDKit Postgres cartridge

- ✓ Direct incorporation of the Tantivy indexing/search engine Rust crate
- ✓ Enables straight-forward multi-threaded indexing and searches
- ✓ Allows easier set-up within our software development stack
- ✓ Main drawback is not having Boost serialization of molecules → we try to avoid unnecessary re-parsing of molecules



Current schema breakdown

SMILES

- Standardized/canonicalized

Descriptors

- Integer fields (e.g. NumAtoms, NumHBD)
- Float fields (e.g. exactmw, TPSA)

Fingerprint

- Stored as bytes field

Miscellaneous metadata

- Stored as nested json field
- Very flexible
- E.g. { "orgs": [1, 2, 3], "amounts_mg": [1, 5, 10], "scaffolds": [10, 12, 13] }




Cheminée indexes scaffolds

- ✓ Currently not using a graph database architecture (potential improvement for the future)
- ✓ Instead, we use an empirical set of 1000 scaffolds to speed up structure searches
 - ❑ Run ScaffoldNetwork on 100K randomized PubChem compounds
 - ❑ Rank scaffolds according to prevalence
- ✓ For every indexed compound, assign scaffolds from the empirical set




Structure searches


Query




Superstructure



Substructure

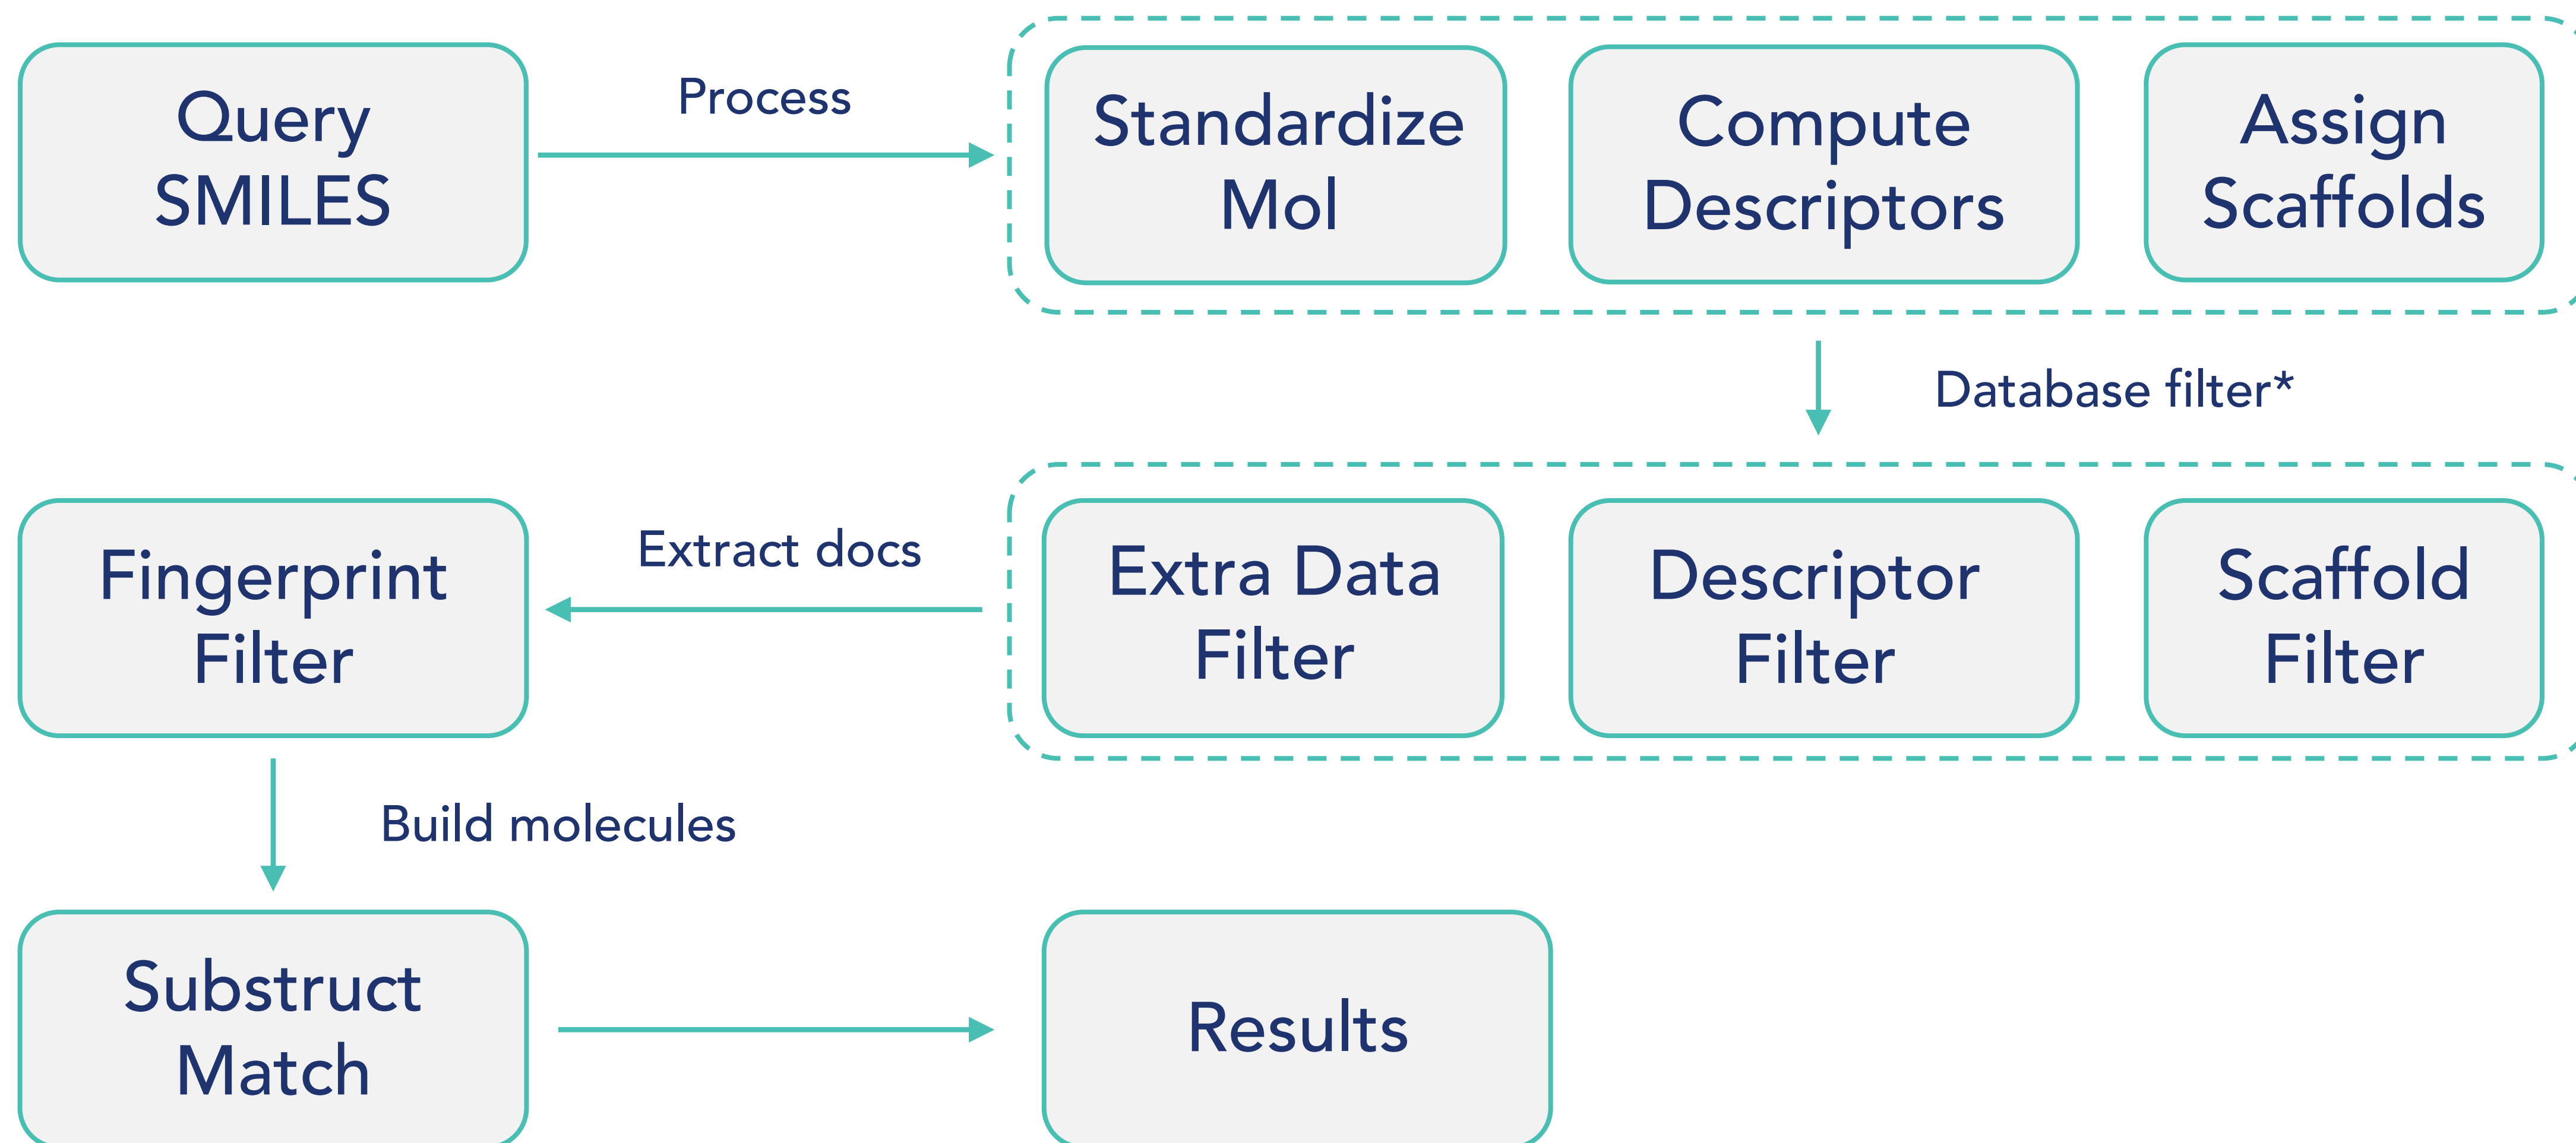


Exact





Substructure search breakdown



*Tautomers processed in parallel



Cheminée can likely get faster

- ✓ Indexing
 - ☐ ~2 million compounds per hour (with automated multithreading)
- ✓ Substructure searches*
 - ☐ ~10-150 milliseconds
- ✓ Superstructure searches*
 - ☐ ~10-150 milliseconds

*Using a test dataset of 5 million compounds



Overview

1. Scientist.com

Company context and structure search requirements

2. The Rust-y RDKit

Low-level Rust bindings for the RDKit C++

3. Cheminée

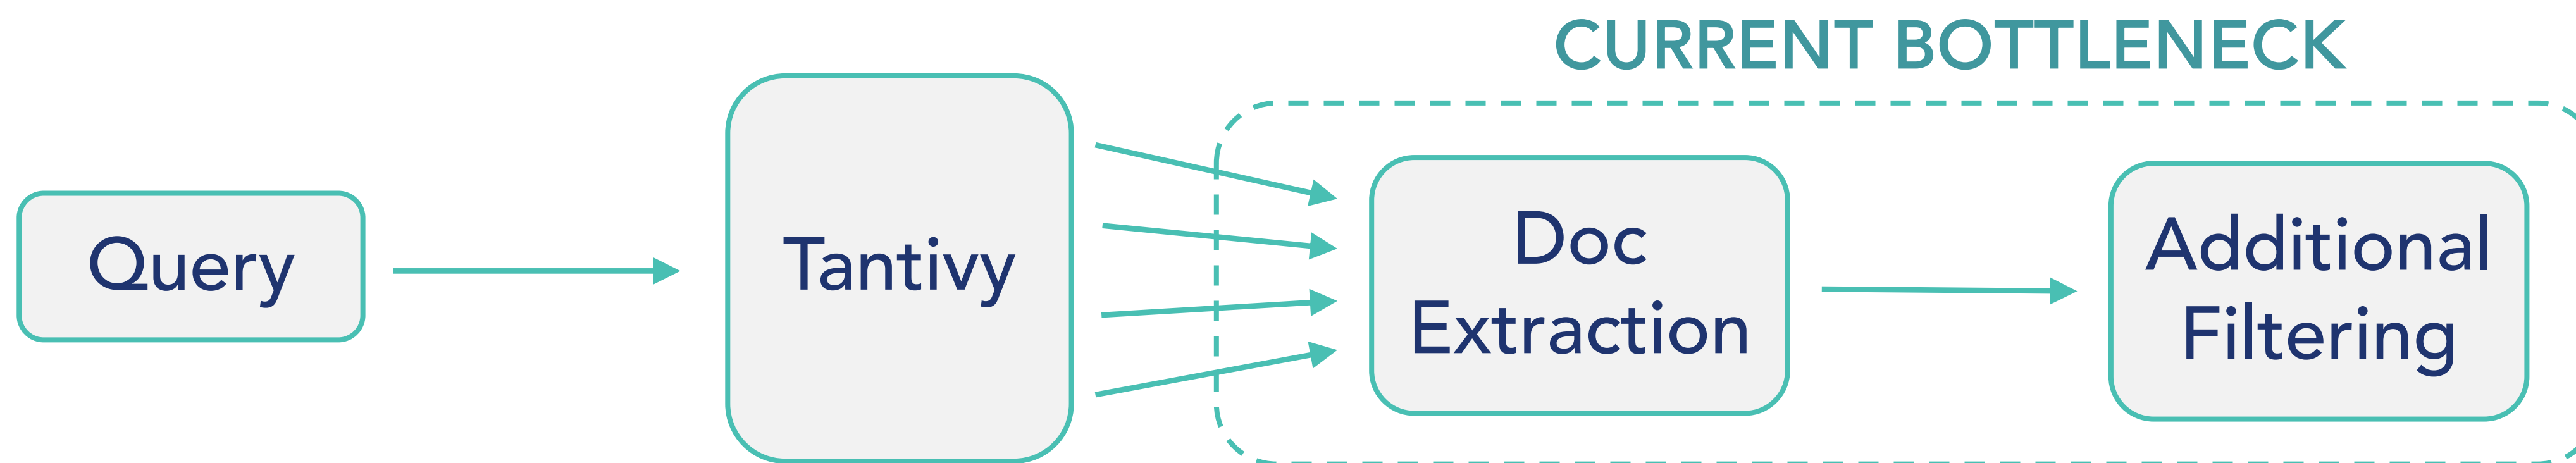
Using the Rust-y RDKit for a stand-alone structure search application

4. Future Directions



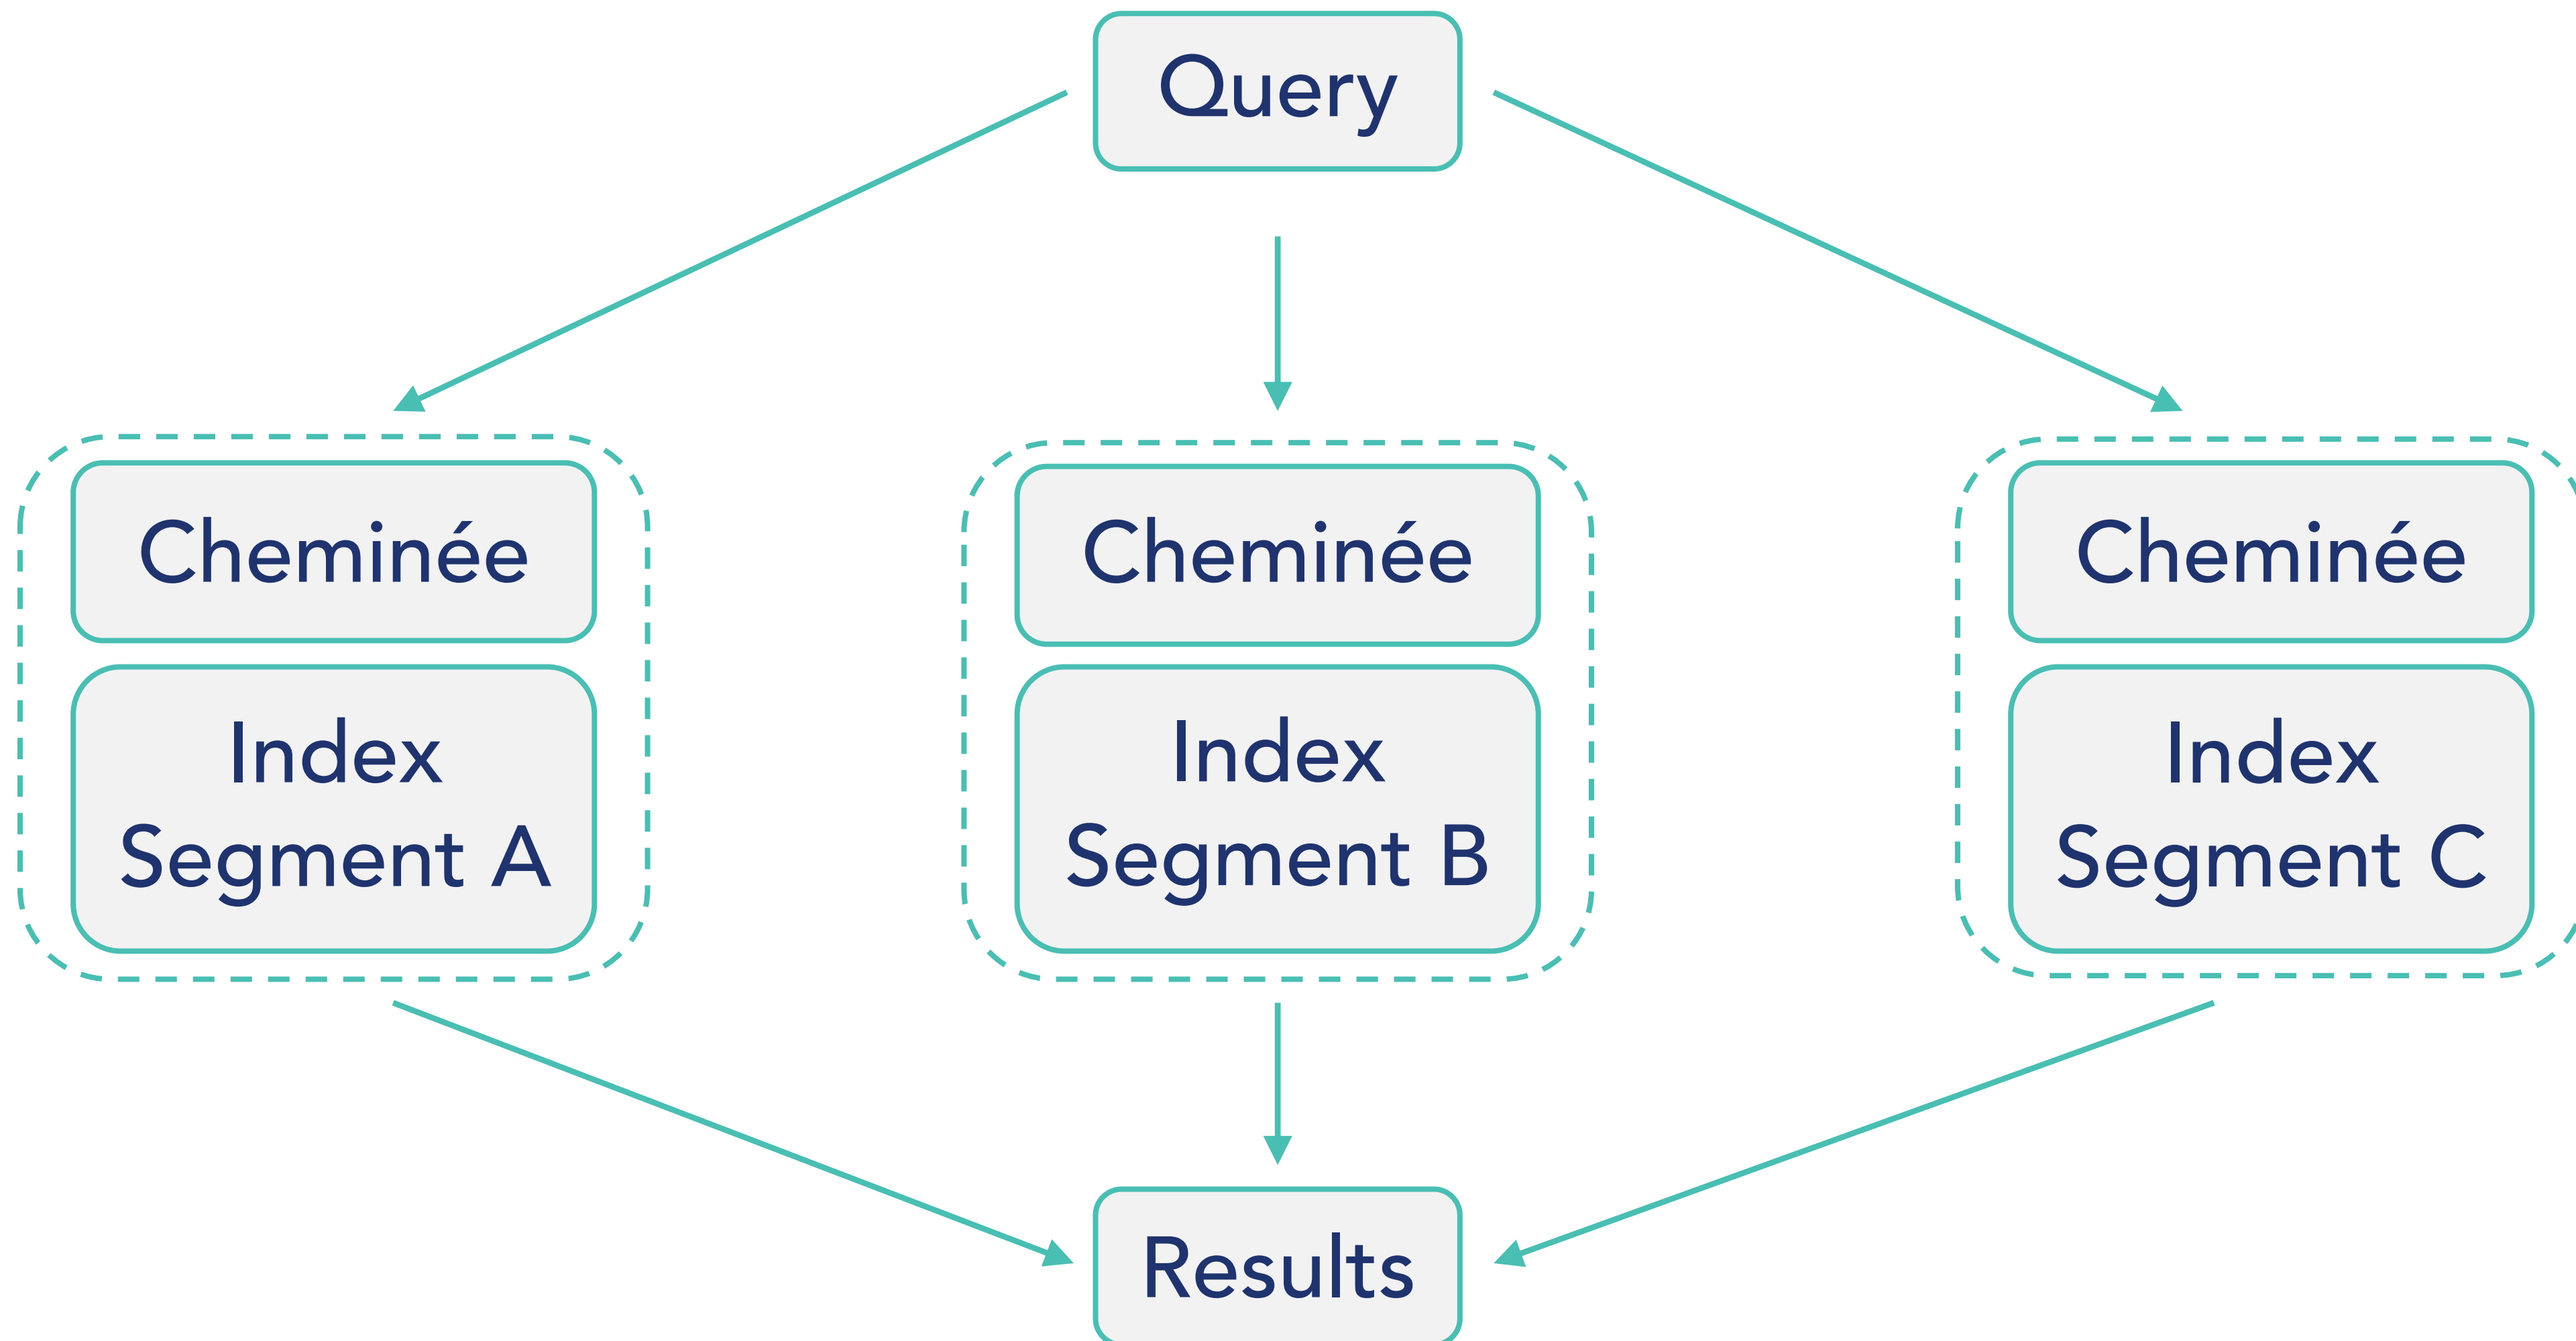
Improvements in the queue

- ✓ Faster indexing
 - ❑ Scaffold assignment involves 1000 substructure matches per compound
 - Probably unnecessary brute force...
- ✓ Faster searching
 - ❑ Push more filtering into Tantivy before doc extraction





Future direction: A horizontally scaled Cheminée*



*Especially for very large libraries



Acknowledgments

- ☆ Xavier Lange
- ☆ Scientist.com
- ☆ RDKit developers
- ☆ RDKit meeting organizers



Contributions are welcome!

The Rust-y RDKit

- ✓ Check out our repo: <https://github.com/rdkit-rs/rdkit>
- ✓ Check out the bindings: <https://docs.rs/rdkit>

Cheminée

- ✓ Check out our repo: <https://github.com/rdkit-rs/cheminee>
- ✓ Spin up our docker container image for some quick testing*

*docker run --rm -dt -p 4001:4001 --name cheminee ghcr.io/rdkit-rs/cheminee:0.1.31